
Policy Optimization with Advantage Decomposition

Anonymous Author(s)

Affiliation

Address

email

Abstract

Policy gradient (PG) on high-dimensional control tasks exhibit its difficulty because of the large variance of the PG estimators. Recent studies show that the Rao-Blackwell theorem (RB) can efficiently reduce the variance of PG by decomposing the action space. However, to invoke RB, previous studies either make strong assumptions or heuristically use the Hessian matrix of the advantage function which is not theoretically sound and also requires specific network architectures to obtain. Inspired by those works, we present algorithmic approaches via pairwise testing and submodular optimization to decompose the action space using the information from the advantage function. Our methods impose weaker assumptions than both previous proposals and are the first ones to provide rigorous performance guarantees. Empirical studies demonstrate the performance improvements on high-dimensional synthetic settings and OpenAI Gym’s MuJoCo continuous control tasks.

1 Introduction

Policy gradient (PG) methods [19] have been widely applied to various challenging problems including video games [12], robotics [8], and continuous control tasks [16, 10]. A major challenge of PG is the high variance of the gradient estimator. Since its inception, the variance of the PG estimator has been continuously improved [19, 17, 12, 6, 11, 4, 18]. Especially, when dealing with high-dimensional action spaces, Rao-Blackwell theorem (RB) [2] based methods [20, 9] demonstrate significant efficiency. To use RB on the PG estimator, the algorithm is required to find the underlying dependency structure of the action space. The structure information is important to keep the unbiasedness of the estimator while finding the estimator with the optimal variance.

Several studies have been proposed. [20] assumes that the policy distribution is fully factorized and decomposes the PG estimator into as many components as the dimension of the action. The action-dependent estimator (ADFB) proposed in [20] also uses averaged action value as the baseline function. However, the fully factorized assumption is relatively strong in practice. Also, there’s an inconsistency: the demonstrated algorithm uses the action value of averaged action as the baseline instead of the averaged action value proposed in its theory. [9] proposes policy optimization with second-order advantage information (POSA) which leverages the block diagonal structure of the advantage Hessian matrix. It successfully finds the non-trivial action space structure, but may suffer from the following issues. First, POSA focuses on the elements of Hessian. Though the Hessian elements corresponding to two independent action dimensions are supposed to be zero, the converse does not hold. Hence, finding the block diagonal structure then does not guarantee the correct partition of the action space. Second, POSA finds the independent partition within the action space using Hessian which indicates pairwise dependency. That is also a necessary condition but not a sufficient one. A simple counter-example would be having three boolean variables that sum up to one: the variables are pairwise independent, but mutually dependent. These inherent weaknesses motivate the rigorous study of the application of RB on PG.

We present policy optimization with advantage decomposition (POAD). Our proposed algorithms make the assumption that the advantage function can be decomposed into multiple components which is weaker than both the assumptions made in POSA and in ADFB. The assumption yields exactly the same estimator which was demonstrated but not explained in ADFB but not explained. Our assumption is also purely based on value learning which allows the algorithm to make the best approximation and leverages the value function information for the policy learning. Finding the decomposition of the action-value function also addresses the flaws in POSA. As the algorithm no longer relies on the pairwise dependencies of the dimensions of the action space, the results are theoretically sound.

We propose two algorithms for computing a decomposition. Both are based on a novel “dependence score” that can be applied to any two disjoint subsets of variables. This score is described by a simple formula that can be estimated to within arbitrary accuracy by random sampling. Both the approaches are efficient (in sublinear time) and have theoretically sound performance guarantees. Our Pairwise Estimates (PE) algorithm computes dependence score estimates for every pair among n variables that correspond to n dimensions of the action space. The estimates are used to build a weighted graph where each node is a variable and the weight of each edge is the corresponding score. The edges are keeping removed in decreasing weight order until only the desired number k of connected components is left. We prove that the output partition is a factor- $O(kn^2)$ approximation of the best possible one. Our Submodular Minimization (SM) algorithm achieves a significantly better almost factor-4 approximation guarantee. The main insight behind this algorithm is that the dependence score is a submodular function of the bipartition. We then reduce the problem to computing the optimal partition, which could be solved efficiently by existing studies of submodular minimization algorithms [5, 14, 7].

We evaluate both our algorithms on a variety of reinforcement learning tasks, including the high-dimension environments from OpenAI Gym’s MuJoCo continuous control tasks. POAD outperforms the existing approaches with both the PE algorithm and the SM algorithm. Especially, POAD with SM has the best overall performance and the best time complexity. Also, POAD with our decomposition algorithms is more stable than POSA’s [9] which uses the Hessian matrix to partition the action space.

2 Preliminaries

2.1 Policy Optimization

We consider policy optimization in the discrete-time Markov decision process (MDP) setting, denoted as the tuple $(\mathcal{S}, \mathcal{A}, \mathcal{T}, r, \rho_0, \gamma)$. That includes $\mathcal{S} \in \mathbb{R}^m$ is the m dimensional state space, $\mathcal{A} \in \mathbb{R}^n$ is the n dimensional action space, $\mathcal{T} : \mathcal{S} \times \mathcal{A} \times \mathcal{S} \rightarrow \mathbb{R}^+$ is the environment transition probability function, $r : \mathcal{S} \times \mathcal{A} \rightarrow \mathbb{R}$ is the reward function, ρ_0 is the initial state distribution and $\gamma \in (0, 1]$ is the unnormalized discount factor. Policy optimization learns a stochastic policy $\pi_\theta : \mathcal{S} \times \mathcal{A} \rightarrow \mathbb{R}_+$, which is parameterized by θ , to maximize the expected cumulative reward

$$J(\theta) = \mathbb{E}_{s \sim \rho_\pi, a \sim \pi} \left[\sum_{t=0}^{\infty} \gamma^t r(s_t, a_t) \right],$$

where $\rho_\pi(s) = \sum_{t=1}^{\infty} \gamma^{t-1} \mathbb{P}(s_t = s)$ is the discounted state visitation distribution. Define the value function

$$V^\pi(s_t) = \mathbb{E}_\pi \left[\sum_{t' \geq t} \gamma^{t'-t} r(s_{t'}, a_{t'}) | s_t, \pi \right]$$

to be the expected return of policy π at state s_t . Notices that the state-action function as

$$Q^\pi(s_t, a_t) = \mathbb{E}_\pi \left[\sum_{t' \geq t} \gamma^{t'-t} r(s_{t'}, a_{t'}) | s_t, a_t, \pi \right]$$

to be the expected return by policy π after taking the action a_t at the state s_t . We make use of $\hat{Q}^\pi(s_t, a_t)$ and $\hat{V}^\pi(s_t)$ to denote the empirical function approximator of $Q^\pi(s_t, a_t)$ and $V^\pi(s_t)$, respectively. Define the advantage function to be the gap between the value function and the action-value, as $A^\pi(s_t, a_t) = Q^\pi(s_t, a_t) - V^\pi(s_t)$. For simplicity, we rewrite it as the time-independent formulation $J(\theta) = \mathbb{E}_{\pi, \rho_\pi} [r(s, a)]$.

Subsequently, we give the formulation of Action Dependent Factorization Baselines (ADFB) estimator [20],

$$\nabla_{\theta} J(\theta)_{ADFG} = \sum_{i=1}^n \mathbb{E}_{\pi(a_i)} [\nabla_{\theta} \log \pi(a_i|s) \{Q^{\pi}(s, a) - Q(s, \bar{a}_i, a_{-i})\}],$$

where a_i is i -th action dimension and denote $a_{-i} = a_{[n]/\{i\}}$, and Action Subspace Dependent Gradient (ASDG) estimator [9],

$$\begin{aligned} \nabla_{\theta} J(\theta)_{ASDG} = & \sum_{j=1}^k \mathbb{E}_{\pi(a_{(j)}|s)} [\nabla_{\theta} \log \pi(a_{(j)}|s) (A^{\pi}(s, a_{(j)}) - c(s, (a_{(j)}, \tilde{a}_{(-j)}))) \\ & - \nabla_{\theta} f_j(\theta, s, \xi) \nabla_{a_{(j)}} c_j(s, a_{(j)})], \end{aligned}$$

where the reparametrization term $\nabla_{\theta} f(\theta, s, \xi) \in \mathbb{R}^{N_{\theta} \times n}$ is divided into k parts as $\nabla_{\theta} f = [\nabla_{\theta} f_1, \dots, \nabla_{\theta} f_k]$ and N_{θ} is the dimension of θ . Note that k is the number of partitions of the action space. Similarly, $a_{(j)}$ is j -th action dimension subset and $a_{(-j)}$ is the complementary action dimension set. Specifically, $\tilde{a}_{(-j)}$ denotes that $a_{(-j)}$ is regarded as a constant in the function. Also note that ADFG is the special case of ASDG.

An important observation is that ASDG is unbiased and with lower variance under the assumption that the advantage is decomposed by the action spaces. We formulate the assumption as below with an error term, and thereafter focus on minimizing the error term in our algorithm.

Assumption 1 (Advantage Decomposition Assumption). *Assume that the advantage function $A^{\pi}(s, a)$ can be locally approximated by decomposition function with respect to a , that is,*

$$A^{\pi}(s, a) = A_1^{\pi}(s, a_{(1)}) + \dots + A_k^{\pi}(s, a_{(k)}) + D(s, a),$$

where $D(s, a)$ is the approximation error.

2.2 Direct Sum Decomposition

As pointed out in the Assumption 1, we consider the general family $F(\mathbf{V})$ as a function of a product set. A direct sum decomposition of F is a partition $(\mathbf{X}_1, \dots, \mathbf{X}_k)$ of the set of variables \mathbf{V} such that

$$F(\mathbf{V}) = F_1(\mathbf{X}_1) + \dots + F_k(\mathbf{X}_k) \quad (1)$$

for some functions F_1, \dots, F_k . We are interested in approximate sublinear-time algorithms for computing direct sum decompositions given oracle access to F .

Given a partition $\mathbf{X}_1, \dots, \mathbf{X}_k$ of the variables, the quality of the decomposition with respect to this partition is measured by

$$\delta^p(\mathbf{X}_1, \dots, \mathbf{X}_k) = \min_{F_1, \dots, F_k} \|F(X_1, \dots, X_k) - F_1(X_1) - \dots - F_k(X_k)\|_p$$

where $\|\cdot\|_p$ is the ℓ^p -norm over the product measure $\|F\| = \mathbb{E}[|F(X_1, \dots, X_k)|^p]^{1/p}$. In Sections 3 and 4 the norm index p can be any integer at least 1. In Section 5 $|\cdot|$ is absolute value, and p equals 2. When p is clear from context we omit it from the notation.

The objective is to find an approximation of the best-possible partition, which minimizes the objective

$$\delta_2^p(F) = \min \delta^p(\mathbf{X}, \overline{\mathbf{X}})$$

where the minimum is taken over all bipartitions $(\mathbf{X}, \overline{\mathbf{X}})$ of \mathbf{V} and $\overline{\mathbf{X}}$ denotes the complement of \mathbf{X} . More generally, we consider partitions into (at least) k nonempty sets, giving rise to the objective value $\delta_k^p(F) = \min \delta^p(\mathbf{X}_1, \dots, \mathbf{X}_k)$.

Our algorithms are based on the following dependence estimator inspired by the rank-1 test of [3]. Let \mathbf{X} and \mathbf{Y} be two disjoint sets of variables. The dependence estimator $D_F(\mathbf{X}, \mathbf{Y})$ is the random variable

$$D_F = F(X, Y, Z) + F(X', Y', Z) - F(X', Y, Z) - F(X, Y', Z)$$

where X, X' are independent samples of the \mathbf{X} variable, Y, Y' are independent samples of the \mathbf{Y} variable, and Z is a random sample of the remaining variables.

117 If F decomposes into a direct sum that partitions the \mathbf{X} and \mathbf{Y} variables then D_F equals zero.
 118 Conversely, $\|D_F\|_p$ measures the quality of the approximation in the ℓ^p -norm.

119 In the analysis it will be convenient to use the notation $F \approx_\delta G$ for $\|F - G\|_p \leq \delta$. The following
 120 two facts are immediate

121 **Fixing:** If $F(X, Z) \approx_\delta G(X, Z)$ then $F(\underline{X}(Z), Z) \approx_\delta G(\underline{X}(Z), Z)$ for some function $\underline{X}(Z)$.

122 **Triangle inequality:** If $F \approx_\delta G$ and $G \approx_{\delta'} H$ then $F \approx_{\delta+\delta'} H$.

123 2.3 Submodularity Optimization

124 In preparation towards our algorithms in Section 5, we recall the definition of a submodular function
 125 and its Lovász extension.

126 **Definition 2** (Submodular set function). *If \mathbf{V} is a finite set, a submodular function is a set function*
 127 *$D : 2^{\mathbf{V}} \rightarrow \mathbb{R}$, where $2^{\mathbf{V}}$ denotes the power set of \mathbf{V} . For every $S, T \subseteq \mathbf{V}$, we have*

$$D(S) + D(T) \geq D(S \cup T) + D(S \cap T).$$

128 **Definition 3** (Lovász extension). *Consider any vector $x = \{x_1, \dots, x_n\}$ such that each $0 \leq x_i \leq 1$.*
 129 *Then the Lovász extension is defined as $D^L = \mathbb{E}(D(\{i | x_i \geq \lambda\}))$ where the expectation is over λ*
 130 *chosen from the uniform distribution on the interval $[0, 1]$.*

131 The Lovász extension is a convex function if and only if the function D is submodular. Therefore, it
 132 reduces submodular minimization of a discrete function to convex minimization of a continuous one.
 133 Several algorithms that rely on this connection have been proposed [5, 14, 7]; see [1] for a survey.

134 3 Estimating the Quality of a Bipartition

135 In this Section we show that $\|D_F(\mathbf{X}, \mathbf{Y})\|_p$ is an approximate estimator for the quality $\delta^p(\mathbf{X}, \mathbf{Y})$ of
 136 a decomposition, namely

$$\delta^p(\mathbf{X}, \mathbf{Y}) \leq \|D_F(\mathbf{X}, \mathbf{Y})\|_p \leq 4 \cdot \delta^p(\mathbf{X}, \mathbf{Y}). \quad (2)$$

137 We then state standard bounds on the quality of the sampling approximators for this value. Taken
 138 together, we obtain efficient factor $(4 + \varepsilon)$ -approximation algorithms for $\delta^p(\mathbf{X}, \mathbf{Y})$ of a given
 139 partition:

140 **Theorem 4.** *There is an algorithm that given a bipartition \mathbf{X}, \mathbf{Y} of the variables and a parameter*
 141 *$\varepsilon > 0$, outputs a value between $\delta^p(\mathbf{X}, \mathbf{Y})$ and $4\delta^p(\mathbf{X}, \mathbf{Y}) + \varepsilon$ with probability at least $1 - \gamma$ from*

$$\min\{K^p \|F\|_{2p}^{2p} / \gamma \varepsilon^{2p}, K^p \|F\|_{\infty}^{2p} \log(1/\gamma) / \varepsilon^{2p}\} \quad (3)$$

142 *queries to F in time linear in the number of queries, where $K > 0$ is a constant.*

143 The first bound has a linear dependence on $1/\gamma$ assuming merely that F is bounded in the $2p$ -norm.
 144 The second bound has a much better, logarithmic dependence on $1/\gamma$ but requires F to be uniformly
 145 bounded. Intermediate bounds should also be possible but for simplicity we only provide the two
 146 extremes.

147 The analysis in fact applies to any pair of disjoint subsets \mathbf{X}, \mathbf{Y} that do not necessarily partition all
 148 the variables. In this more general setting (which will be useful in Section 4) distance is measured by
 149 the formula

$$\delta^p(\mathbf{X}, \mathbf{Y}) = \min_{A, B} \|F(X, Y, Z) - A(X, Z) - B(Y, Z)\|_p.$$

150 **Claim 5** (Completeness of D_F). *For all disjoint \mathbf{X}, \mathbf{Y} , $\|D_F(\mathbf{X}, \mathbf{Y})\|_p \leq 4 \cdot \delta^p(\mathbf{X}, \mathbf{Y})$.*

151 *Proof.* By definition of $\delta(\mathbf{X}, \mathbf{Y})$ there exists a decomposition of the form

$$F(V) = A(X, Z) + B(Y, Z) + D(X, Y, Z),$$

152 where $\|D(X, Y, Z)\| = \delta(\mathbf{X}, \mathbf{Y})$. In the expansion of D_F all the A and B terms cancel out, leaving

$$\begin{aligned} \|D_F(\mathbf{X}, \mathbf{Y})\| &= \|D(X, Y, Z) + D(X', Y', Z) - D(X, Y', Z) - D(X', Y, Z)\| \\ &\leq \|D(X, Y, Z)\| + \|D(X', Y', Z)\| + \|D(X, Y', Z)\| + \|D(X', Y, Z)\| \\ &= 4\delta(\mathbf{X}, \mathbf{Y}). \end{aligned} \quad \square$$

153 **Claim 6** (Soundness of D_F). *For all disjoint \mathbf{X}, \mathbf{Y} , $\delta^p(\mathbf{X}, \mathbf{Y}) \leq \|D_F(\mathbf{X}, \mathbf{Y})\|_p$* ¹

154 *Proof.* Let $\varepsilon = \|D_F(\mathbf{X}, \mathbf{Y})\|$. Then

$$F(X, Y, Z) \approx_\varepsilon F(X, Y', Z) - F(X', Y, Z) - F(X', Y', Z).$$

155 By fixing, there exist functions $x' = \underline{X}'(Z)$ and $y' = \underline{Y}'(Z)$ for which

$$F(X, Y, Z) \approx_\varepsilon F(\underline{X}'(Z), \underline{Y}'(Z), Z) - F(X, \underline{Y}'(Z), Z) - F(\underline{X}'(Z), Y, Z).$$

156 Therefore $F(X, Y, Z) \approx_e A(X, Z) + B(Y, Z)$, where $A(X, Z) = F(X, \underline{Y}'(Z), Z)$ and
 157 $B(Y, Z) = F(\underline{X}'(Z), Y, Z) - F(\underline{X}'(Z), \underline{Y}'(Z), Z)$. \square

158 To finish the proof of Theorem 4 we need the following (standard) sampling bounds derived from
 159 Chebyshev's and Chernoff's inequalities.

160 **Claim 7.** *The empirical average of $|f(X)|^p$ over N random samples X is within ε^p of $\|f\|_p^p$ except
 161 with probability $\min\{\|f\|_{2p}^{2p}/N\varepsilon^{2p}, 2\exp(-N\varepsilon^{2p}/2\|f\|_\infty^{2p})\}$.*

162 **Claim 8.** *Assume $t, \hat{t} \geq 0$. If $t^p \leq \hat{t}^p \leq t^p + (\varepsilon/2)^p$ then $t \leq \hat{t} \leq t + \varepsilon$.*

163 *Proof.* The left-hand inequalities are immediate. For the right-hand ones we start we consider two
 164 cases. If $t \leq \varepsilon/2$, then $\hat{t}^p \leq 2(\varepsilon/2)^p \leq \varepsilon \leq t + \varepsilon$. If $t > \varepsilon/2$ then

$$\hat{t} - t \leq \frac{\hat{t}^p - t^p}{t^{p-1}} \leq \frac{(\varepsilon/2)^p}{(\varepsilon/2)^{p-1}} \leq \varepsilon. \quad \square$$

165 **Claim 9.** *The value $\|D_F(\mathbf{X}, \mathbf{Y})\|_p$ can be estimated within ε from (3) queries to F in linear time
 166 with probability $1 - \gamma$.*

167 *Proof.* By the triangle inequality, $\|D_F(\mathbf{X}, \mathbf{Y})\|_p \leq 4\|F\|_p$. By convexity $\|D_F(\mathbf{X}, \mathbf{Y})\|_p^p \leq$
 168 $4^{p-1}\|F\|_p^p$. By Claim 7, an $(1 - \gamma)$ -probability estimate that lies between $\|D_F(\mathbf{X}, \mathbf{Y})\|_p^p$ and
 169 $\|D_F(\mathbf{X}, \mathbf{Y})\|_p^p + (\varepsilon/2)^p$ can be computed by taking the empirical average of (3) samples of D_F .
 170 Each sample of D_F can be obtained using four queries to F . The quality of the estimate then follows
 171 from Claim 8. \square

172 Theorem 4 now follows from Claim 9 and inequality (2).

173 4 Approximate Partitioning via Pairwise Estimates

174 We analyze the following algorithm for computing an approximate decomposition.

Algorithm 1: Direct sum decomposition algorithm

Input: number of partitions k .

Output: partition \mathcal{P} .

for each iteration do

 (Step 1) For every pair of distinct variables \mathbf{x}, \mathbf{y} , calculate an estimate $\hat{e}(\mathbf{x}, \mathbf{y})$ for

$e(\mathbf{x}, \mathbf{y}) = \|D_F(\{\mathbf{x}\}, \{\mathbf{y}\})\|_p$;

 (Step 2) Create a weighted graph with vertices \mathbf{V} and weights $\hat{e}(\mathbf{x}, \mathbf{y})$. Order the edges in
 increasing weight and keep removing them until the graph partitions into k components.²;

end

175 **Theorem 10.** *Assuming $e(\mathbf{x}, \mathbf{y}) \leq \hat{e}(\mathbf{x}, \mathbf{y}) \leq e(\mathbf{x}, \mathbf{y}) + \varepsilon$ for all \mathbf{x} and \mathbf{y} ,*

$$\delta^p(\mathcal{P}) \leq (8k - 10)n^2(4\delta_k^p(F) + \varepsilon). \quad (4)$$

176 **Corollary 11.** *If Step 1 of the algorithm is implemented by empirically averaging over*

$$\min\{K^p\|F\|_{2p}^{2p}/\gamma\varepsilon^{2p}, K^p\|F\|_\infty^{2p}\log(n/\gamma)/\varepsilon^{2p}\}$$

177 *random samples then approximation guarantee (4) holds with probability at least $1 - \gamma$. Note that K
 178 *is an absolute constant.**

¹In the case of Boolean functions under uniform measure this was proved by David et al. [3].

5 Approximate Partitioning via Submodular Function Minimization

Let $f(\mathbf{X}) = \|D_F(\mathbf{X}, \overline{\mathbf{X}})\|_2^2 = \mathbb{E}[D_F(\mathbf{X}, \overline{\mathbf{X}})^2]$. The main result of this section is

Theorem 12. f is a symmetric submodular function.

By (2), $\delta_2^2(F) \leq \min \|D_F(\mathbf{X}, \overline{\mathbf{X}})\|_2 \leq 4\delta_2^2(F)$, where the minimum is taken over all nontrivial bipartitions. Since submodular functions can be minimized efficiently [5, 13, 14, 7], by Theorem 12 it is possible to efficiently compute a factor-4 approximation to $\delta_2^2(F)$ given oracle access to f .

Corollary 13. *There is an algorithm that, given oracle access to F , runs in time polynomial in n , $1/\varepsilon$, and $\|F\|_4/\gamma$ or $\|F\|_\infty \log(1/\gamma)$, and outputs a set \mathbf{X} such that $\delta_2(F) \leq f(\mathbf{X}) \leq 4\delta_2(F) + \varepsilon$ with probability at least $1 - \gamma$.*

In the remaining part of this section we prove Theorem 12.

Claim 14. For every partition (\mathbf{X}, \mathbf{Y}) , $\mathbb{E}[D_F(\mathbf{X}, \mathbf{Y})^2] = 4 \cdot \mathbb{E}[F(\mathbf{X}, \mathbf{Y}) \cdot D_F(\mathbf{X}, \mathbf{Y}, \mathbf{X}', \mathbf{Y}')]]$.

Proof. Let X_0, X_1 be independent copies of \mathbf{X} and Y_0, Y_1 be independent copies of \mathbf{Y} . We can express D_F in the form

$$D_F = \sum_{i,j \in \{0,1\}} (-1)^{i+j} F(X_i, Y_j).$$

Then

$$\begin{aligned} \mathbb{E}[D_F^2] &= \mathbb{E} \sum_{i,j \in \{0,1\}} (-1)^{i+j} F(X_i, Y_j) \cdot \sum_{i',j' \in \{0,1\}} (-1)^{i'+j'} F(X_{i'}, Y_{j'}) \\ &= \mathbb{E} \sum_{i,j,i',j' \in \{0,1\}} (-1)^{i+j+i'+j'} F(X_i, Y_j) \cdot F(X_{i'}, Y_{j'}) \\ &= \sum_{i,i',j,j' \in \{0,1\}} (-1)^{(i \oplus i') + (j \oplus j')} \mathbb{E}[F(X_i, Y_j) \cdot F(X_{i'}, Y_{j'})], \end{aligned}$$

where \oplus is addition modulo 2. By symmetry of the variables

$$\mathbb{E}[F(X_i, Y_j) \cdot F(X_{i'}, Y_{j'})] = \mathbb{E}[F(X_0, Y_0) \cdot F(X_{i \oplus i'}, Y_{j \oplus j'})]$$

substituting in the previous expression and changing the order of the expectation and the summation we obtain

$$\mathbb{E}[D_F^2] = \mathbb{E} F(X_0, Y_0) \sum_{i,i',j,j' \in \{0,1\}} (-1)^{(i \oplus i') + (j \oplus j')} F(X_{i \oplus i'}, Y_{j \oplus j'}).$$

The summation is exactly $4D_F(X_0, Y_0, X_1, Y_1)$ as desired. \square

Let $D_F(\mathbf{X}, \mathbf{Y} \mid \mathbf{Z})$ denote $D_{F_Z}(\mathbf{X}, \mathbf{Y})$, where F_Z is obtained from F by fixing \mathbf{Z} to \mathbf{Z} .

Claim 15. $D_F(\mathbf{X}\mathbf{Y}, \mathbf{Z}\mathbf{W}) + D_F(\mathbf{X}\mathbf{Z}, \mathbf{Y}\mathbf{W}) - D_F(\mathbf{X}\mathbf{Y}\mathbf{Z}, \mathbf{W}) - D_F(\mathbf{X}, \mathbf{Y}\mathbf{Z}\mathbf{W})$ is equal to $D_F(\mathbf{Y}, \mathbf{Z} \mid \mathbf{X}\mathbf{W}') + D_F(\mathbf{Y}, \mathbf{Z} \mid \mathbf{X}'\mathbf{W})$.

Proof. Expanding the left-hand side yields sixteen terms of the form $F(X^?Y^?Z^?W^?)$ with a leading plus or minus sign, where the question mark indicates a primed or unprimed superscript. Each of the terms $F(XY\mathbf{Z}\mathbf{W})$ and $F(X'Y'Z'W')$ occurs exactly twice with a plus sign and twice with a minus sign, so those eight terms cancel out. There remain four terms of the type $F(XY^?Z^?W')$ and four of the type $F(X'Y^?Z^?W)$. Among those, the terms where exactly one of Y, Z is primed are negative and the others are positive. Therefore the type $F(XY^?Z^?W')$ terms add up to $D_F(\mathbf{Y}, \mathbf{Z} \mid \mathbf{X}\mathbf{W}')$ and the type $F(X'Y^?Z^?W)$ terms add up to $D_F(\mathbf{Y}, \mathbf{Z} \mid \mathbf{X}'\mathbf{W})$. \square

Claim 16. For independent Y, Z , $\mathbb{E}_Y \text{Var}_Z F(Y, Z) \geq \text{Var}_Z \mathbb{E}_Y F(Y, Z)$.

Proof. Set $F' = F - \mathbb{E}_Z[F]$. Then

$$\mathbb{E}_Y \text{Var}_Z F = \mathbb{E}_Y \mathbb{E}_Z [F'^2] = \mathbb{E}_Z [\mathbb{E}_Y [F'^2]],$$

while

$$\text{Var}_Z \mathbb{E}_Y F = \mathbb{E}_Z [(\mathbb{E}_Y F - \mathbb{E}_Z \mathbb{E}_Y F)^2] = \mathbb{E}_Z [\mathbb{E}_Y [F']^2].$$

The fact follows from the Cauchy-Schwarz inequality. \square

211 **Claim 17.** $\mathbb{E}[F(XYZW)D_F(YZY'Z' | X'W)]$ is nonnegative.

212 *Proof.* Without loss of generality we may assume W is empty. Then the expression is the difference
213 of the terms

$$\begin{aligned} ev &= \mathbb{E}[F(XYZ)F(X'YZ)] - \mathbb{E}[F(XYZ)F(X'YZ')] \quad \text{and} \\ ve &= \mathbb{E}[F(XYZ)F(X'Y'Z)] - \mathbb{E}[F(XYZ)F(X'Y'Z')]. \end{aligned}$$

214 Let $\hat{F}(YZ) = \mathbb{E}_X[F(XYZ)]$. Then

$$ev = \mathbb{E}_{YZ}[\hat{F}^2] - \mathbb{E}_Y[\mathbb{E}_Z[\hat{F}^2]] = \mathbb{E}_Y \text{Var}_Z[\hat{F}].$$

215 and

$$ve = \mathbb{E}_Z[\mathbb{E}_Y[\hat{F}^2]] - \mathbb{E}_{ZY}[\hat{F}^2] = \text{Var}_Z \mathbb{E}_Y[\hat{F}].$$

216 The claim now follows from Fact 16. □

217 *Proof of Lemma 12.* Proving submodularity of f amounts to verifying the inequality

$$\mathbb{E}[D_F(\mathbf{X}\mathbf{Y}, \mathbf{Z}\mathbf{W})^2] + \mathbb{E}[D_F(\mathbf{X}\mathbf{Z}, \mathbf{Y}\mathbf{W})^2] - \mathbb{E}[D_F(\mathbf{X}\mathbf{Y}\mathbf{Z}, \mathbf{W})^2] - \mathbb{E}[D_F(\mathbf{X}, \mathbf{Y}\mathbf{Z}\mathbf{W})^2] \geq 0.$$

218 for every partition $(\mathbf{X}, \mathbf{Y}, \mathbf{Z}, \mathbf{W})$ of the variables. This follows by applying Claims 14, 15, and 17
219 to each of the two expressions in order. □

220 6 Experiments

221 6.1 Implementation Details

222 Our setting is similar to the one used in [9]³. Our algorithm uses PPO [16] together with GAE
223 [15]. We use three neural networks including a policy network for PPO and a value network for
224 GAE. Both of the networks have the same architecture as is in the original A2C and PPO work
225 [12, 16]. We utilize a third network which estimates the advantage function $A^\mu(s, a)$. Unlike [9]
226 where $\nabla_{aa} A^\mu(s, a)$ has to be computed efficiently, our advantage network does not require any
227 special neural network architecture to achieve that. In our implementation, $A^\mu(s, a)$ first encodes
228 the state s with a hidden layer and concatenates it with the action a . It then feeds the output into a
229 two-layer MLPs. Other parameters are consistent with those in [9].

230 **Pairwise Estimation Algorithm** Implementation-wise, we use an alternative version of the algorithm
231 which shares the same theoretical guarantee but empirically runs faster. Instead of adding all the
232 edges at the beginning and remove one at a step, we keep adding edges in decreasing weight until just
233 before the connected component count drops below k . During the pairwise estimation, we collect the
234 batch data and sample a second action for each of the state (apart from the action that is interacting
235 with the environment and generating the trajectories). The second action is used as the testing term
236 X' and Y' . We have the pairwise dependence scores (weights of the edge) averaged within each of
237 the batches.

238 **Submodular Minimization Algorithm** While the [5] algorithm is considered impractical, we im-
239 plement the submodular minimization with enumerating the possible partitions. We do not know if
240 any other efficient submodular minimization algorithm (see [1, Chapter 10] for a survey) provides
241 similar optimality guarantees when provided with an approximate oracle as is in our settings and we
242 plan to investigate this in the future.

243 **Synthetic High-Dimensional Action Spaces** Similar to previous works, we design a synthetic
244 environment with a one-step MDP where the reward is quadratic with respect to a and the quadratic
245 term is block-diagonal. As the environment setting satisfies Assumption 1, it achieves almost-optimal
246 performance which is similar to POSA.

247 6.2 OpenAI Gym’s MuJoCo Environments

248 Fig. 1 shows the results on several of the high-dimensional continuous control tasks from OpenAI
249 Gym’s MuJoCo Environments. The legend POAD_PE and POAD_SM stand for our proposed

³We will publish our code upon the acceptance of this paper.

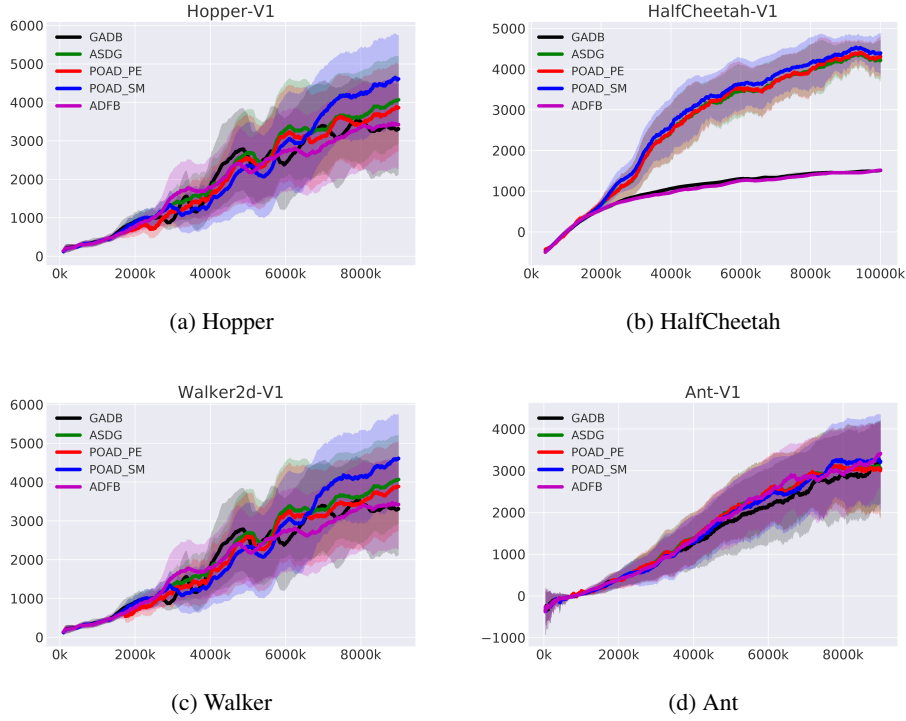


Figure 1: Comparison between three baselines (ADFB, GADB, ASDG) and our two direct sum testing algorithms (POAD_PE and POAD_SM) on various OpenAI Gym Mujoco continuous control tasks, including Hopper-V1, HalfCheetah-V1 Walker2d-V1 and Ant-V1 . Notice that POAD_PE is the pairwise estimation algorithm and POAD_SM is the submodular Minimization.

algorithms. POAD_SM has the best overall performance which is consistent across various of the settings. Our algorithm based on pairwise estimates has similar performance compared with previous Hessian-based algorithm POSA.

We observe that RB works exceptionally well on HalfCheetah, which coincides with the fact that the approximation error is relatively low in such environment. As an intuition, the policy optimization algorithm has to balance accuracy (i.e., the satisfaction of its model assumption) and efficiency (i.e., the variance reduction performance) in practice. The rigorous analysis of the approximation error in our study helps the understanding of the policy optimization algorithms and their assumptions.

References

- [1] Francis Bach. Learning with submodular functions: A convex optimization perspective. *Foundations and Trends in Machine Learning*, 6(2-3):145–373, 2013.
- [2] George Casella and Christian P Robert. Rao-blackwellisation of sampling schemes. *Biometrika*, 83(1):81–94, 1996.
- [3] Roee David, Irit Dinur, Elazar Goldenberg, Guy Kindler, and Igor Shinkar. Direct sum testing. *SIAM Journal on Computing*, 46(4):1336–1369, 2017.
- [4] Will Grathwohl, Dami Choi, Yuhuai Wu, Geoff Roeder, and David Duvenaud. Backpropagation through the void: Optimizing control variates for black-box gradient estimation. In *International Conference on Learning Representations*, 2018.
- [5] Martin Grötschel, László Lovász, and Alexander Schrijver. The ellipsoid method and its consequences in combinatorial optimization. *Combinatorica*, 1(2):169–197, 1981.

- 270 [6] Shixiang Gu, Timothy Lillicrap, Zoubin Ghahramani, Richard E Turner, and Sergey
271 Levine. Q-prop: Sample-efficient policy gradient with an off-policy critic. *arXiv preprint*
272 *arXiv:1611.02247*, 2016.
- 273 [7] Satoru Iwata, Lisa Fleischer, and Satoru Fujishige. A combinatorial strongly polynomial
274 algorithm for minimizing submodular functions. *J. ACM*, 48(4):761–777, July 2001.
- 275 [8] Sergey Levine, Chelsea Finn, Trevor Darrell, and Pieter Abbeel. End-to-end training of deep
276 visuomotor policies. *Journal of Machine Learning Research*, 17(39):1–40, 2016.
- 277 [9] Jiajin Li and Baoxiang Wang. Policy optimization with second-order advantage information.
278 *arXiv preprint arXiv:1805.03586*, 2018.
- 279 [10] Timothy P Lillicrap, Jonathan J Hunt, Alexander Pritzel, Nicolas Heess, Tom Erez, Yuval Tassa,
280 David Silver, and Daan Wierstra. Continuous control with deep reinforcement learning. *arXiv*
281 *preprint arXiv:1509.02971*, 2015.
- 282 [11] Hao Liu, Yihao Feng, Yi Mao, Dengyong Zhou, Jian Peng, and Qiang Liu. Action-dependent
283 control variates for policy optimization via stein identity. In *International Conference on*
284 *Learning Representations*, 2018.
- 285 [12] Volodymyr Mnih, Adria Puigdomenech Badia, Mehdi Mirza, Alex Graves, Timothy P Lill-
286 icrap, Tim Harley, David Silver, and Koray Kavukcuoglu. Asynchronous methods for deep
287 reinforcement learning. In *International Conference on Machine Learning*, 2016.
- 288 [13] Maurice Queyranne. Minimizing symmetric submodular functions. *Mathematical Programming*,
289 82(1-2):3–12, 1998.
- 290 [14] Alexander Schrijver. A combinatorial algorithm minimizing submodular functions in strongly
291 polynomial time. *J. Comb. Theory, Ser. B*, 80(2):346–355, 2000.
- 292 [15] John Schulman, Philipp Moritz, Sergey Levine, Michael Jordan, and Pieter Abbeel. High-
293 dimensional continuous control using generalized advantage estimation. *arXiv preprint*
294 *arXiv:1506.02438*, 2015.
- 295 [16] John Schulman, Filip Wolski, Prafulla Dhariwal, Alec Radford, and Oleg Klimov. Proximal
296 policy optimization algorithms. *arXiv preprint arXiv:1707.06347*, 2017.
- 297 [17] Richard S Sutton, David A McAllester, Satinder P Singh, and Yishay Mansour. Policy gradient
298 methods for reinforcement learning with function approximation. In *Advances in neural*
299 *information processing systems*, pages 1057–1063, 2000.
- 300 [18] George Tucker, Surya Bhupatiraju, Shixiang Gu, Richard E Turner, Zoubin Ghahramani, and
301 Sergey Levine. The mirage of action-dependent baselines in reinforcement learning. *arXiv*
302 *preprint arXiv:1802.10031*, 2018.
- 303 [19] Ronald J Williams. Simple statistical gradient-following algorithms for connectionist reinforce-
304 ment learning. *Machine learning*, 8(3-4):229–256, 1992.
- 305 [20] Cathy Wu, Aravind Rajeswaran, Yan Duan, Vikash Kumar, Alexandre M. Bayen, Sham Kakade,
306 Igor Mordatch, and Pieter Abbeel. Variance reduction for policy gradient with action-dependent
307 factorized baselines. In *International Conference on Learning Representations*, 2018.

A Proofs of the Pairwise Estimates Algorithm

For a partition \mathcal{P} of the variables, let $\Delta(\mathcal{P}) = \sum \delta(\{x\}, \{y\})$, where the sum is taken over all pairs that cross the partition. We will deduce Theorem 10 from the following bound on $\delta(\mathcal{P})$

Claim 18. *For every k -partition \mathcal{P} , $\delta(\mathcal{P}) \leq (16k - 20)\Delta(\mathcal{P})$.*

The following fact is immediate from the definitions of δ .

Fact 19. *For any partition (U, \bar{U}) such that $X \subseteq U$ and $Y \subseteq \bar{U}$, $\delta(X, Y) \leq \delta(U, \bar{U})$.*

Proof of Theorem 10. By Claim 5 and Fact 19, all edges (x, y) in the optimal partition must satisfy $e(x, y) \leq 4\delta_2(F)$. By our assumption on the quality of the approximations,

$$\hat{e}(x, y) \leq 4\delta_2(F) + \varepsilon. \quad (5)$$

Since the algorithm removes edges in increasing order of weight, all the edges that cross the output partition \mathcal{P} must also satisfy this inequality. Then

$$\begin{aligned} \delta(\mathcal{P}) &\leq (16k - 20)\Delta(\mathcal{P}) && \text{by Claim 18,} \\ &\leq (16k - 20) \sum_{x, y \text{ cross } \mathcal{P}} e(x, y) && \text{by Claim 6,} \\ &\leq (16k - 20) \sum_{x, y \text{ cross } \mathcal{P}} \hat{e}(x, y) && \text{by quality of approximation,} \\ &\leq (16k - 20) \sum_{x, y \text{ cross } \mathcal{P}} 4\delta_2(F) + \varepsilon && \text{by (5),} \\ &\leq (8k - 10)n^2 \cdot (4\delta_2(F) + \varepsilon). \end{aligned}$$

The last inequality holds because there are at most $\binom{n}{2} \leq n^2/2$ pairs of variables crossing the partition. \square

In the case $k = 2$, the leading constant $8k - 10 = 6$ can be improved to 1 by using Claim 21 instead of Claim 18 and the fact that at most $n^2/4$ pairs cross the partition.

It remains to prove Claim 18. We use XX' to denote the union of the variable sets X and X' .

Claim 20. *For disjoint sets of variables X, X', Y , $\delta(XX', Y) \leq \delta(X, Y) + 2\delta(X', Y)$.*

Proof of Claim 20. For simplicity of notation we omit the dependences on Z . Assume that

$$\begin{aligned} F(X, X', Y) &\approx_{\delta} A(X, X') + B(X', Y) \quad \text{and} \\ F(X, X', Y) &\approx_{\delta'} A'(X, X') + B'(X, Y). \end{aligned}$$

By the triangle inequality,

$$A(X, X') + B(X', Y) \approx_{\delta+\delta'} A'(X, X') + B'(X, Y).$$

Fix $X'(Z) = \underline{X}'(Z)$. Writing $C(X) = A(X, \underline{X}') - A'(X, \underline{X}')$ and $D(Y) = B(\underline{X}', Y) - B'(\underline{X}', Y)$ we get that

$$B'(X, Y) \approx_{\delta+\delta'} C(X) + D(Y).$$

By the triangle inequality (with the second equation), we get that

$$F(X, X', Y) \approx_{\delta+2\delta'} A'(X, X') + C(X) + D(Y). \quad \square$$

Claim 21. *For every bipartition X, \bar{X} of the variables, $\delta(X, \bar{X}) \leq 4 \cdot \Delta(X, \bar{X})$.*

Proof. By Claim 20,

$$\delta(X' \{x\}, \{y\}) \leq \delta(X', \{y\}) + 2\delta(\{x\}, \{y\})$$

for all $X' \subseteq X \setminus \{x\}$ and y . Applying this inequality iteratively we conclude that $\delta(X, \{y\}) \leq 2 \sum_{x \in X} \delta(\{x\}, \{y\})$. Also by Claim 20

$$\delta(X, Y' \{y\}) \leq \delta(X, Y') + 2\delta(X, Y' \{y\}),$$

so $\delta(X, Y) \leq 2 \sum_{y \in Y} \delta(X, \{y\})$. Combining the two inequalities we obtain the desired conclusion. \square

335 To extend the proof to larger k , we generalize the first inequality in this sequence to k -partitions.
 336 The same sequence of inequalities then yields the conclusion of Theorem 10. It remains to prove
 337 Claim 18.

338 **Claim 22.** For every $2k$ -partition $(\mathbf{Y}_1, \dots, \mathbf{Y}_k, \mathbf{Z}_1, \dots, \mathbf{Z}_k)$,

$$\delta(\mathbf{Y}_1, \dots, \mathbf{Y}_k, \mathbf{Z}_1, \dots, \mathbf{Z}_k) \leq 2\delta(\mathbf{Y}_1 \mathbf{Z}_1, \dots, \mathbf{Y}_k \mathbf{Z}_k) + 3\delta(\mathbf{Y}_1 \dots \mathbf{Y}_k, \mathbf{Z}_1 \dots \mathbf{Z}_k).$$

339 *Proof.* Assume that

$$\begin{aligned} F(\mathcal{V}) &\approx_{\delta} F_1(Y_1, Z_1) + \dots + F_t(Y_t, Z_t) \\ F(\mathcal{V}) &\approx_{\delta'} A(Y_1, \dots, Y_t) + B(Z_1, \dots, Z_t). \end{aligned}$$

340 By the triangle inequality

$$A(Y_1, \dots, Y_t) + B(Z_1, \dots, Z_t) \approx_{\delta+\delta'} F_1(Y_1, Z_1) + \dots + F_t(Y_t, Z_t).$$

341 Fixing Z_1, \dots, Z_t to values $\underline{Z}_1, \dots, \underline{Z}_t$ we get the decomposition

$$A(Y_1, \dots, Y_t) \approx_{\delta+\delta'} F_1(Y_1, \underline{Z}_1) + \dots + F_t(Y_t, \underline{Z}_t) - B(\underline{Z}_1, \dots, \underline{Z}_t).$$

342 and similarly

$$B(Z_1, \dots, Z_t) \approx_{\delta+\delta'} F_1(\underline{Y}_1, Z_1) + \dots + F_t(\underline{Y}_t, Z_t) - A(\underline{Y}_1, \dots, \underline{Y}_t).$$

343 Plugging these into the second equation gives the desired decomposition. \square

344 *Proof of Claim 18.* We assume that k is a power of two and prove by induction that $\delta(\mathcal{P}) \leq c_k \Delta(\mathcal{P})$,
 345 where c_k is the sequence $c_{2k} = 2c_k + 12$, $c_2 = 4$. The base case $k = 2$ follows from Claim 21.
 346 Assume the claim holds for k and apply Claim 22 to \mathcal{P} . By inductive assumption and Claim 21,

$$\delta(\mathcal{P}) \leq 2 \cdot c_k \Delta(\mathbf{Y}_1 \mathbf{Z}_1, \dots, \mathbf{Y}_k \mathbf{Z}_k) + 3 \cdot 4\Delta(\mathbf{Y}_1 \dots \mathbf{Y}_k, \mathbf{Z}_1 \dots \mathbf{Z}_k).$$

347 Since \mathcal{P} is a refinement of both these partitions, it follows that $\delta(\mathcal{P}) \leq (2c_k + 12)\Delta(\mathcal{P}) = c_{2k}\Delta(\mathcal{P})$,
 348 concluding the induction.

349 The recurrence solves to $c_k = 8k - 12$, proving the claim when k is a power of two. When it is not,
 350 the same reasoning applies to the closest power of two exceeding k (by taking some of the sets in the
 351 partition to be empty), which is at most $2k - 1$, proving the desired bound. \square

352 B Submodular Minimization with Inexact Oracle

353 For a general function F , it is impossible to obtain the exact oracle access to $f = \|D_F\|_2^2$, since
 354 it requires exact evaluation of the expectation $\mathbb{E}[D_F(\mathbf{X}, \overline{\mathbf{X}})^2]$. However, we can obtain arbitrarily
 355 good approximations by sampling. That is, given an input \mathbf{X} for f , accuracy parameters ε and
 356 γ , it can output a value between $f(\mathbf{X})$ and $f(\mathbf{X}) + \varepsilon$ with high probability at least $1 - \gamma$. If the
 357 approximation can be computed in time polynomial in $1/\gamma$ and $1/\varepsilon$, we can say that the approximate
 358 oracle is *efficient*. By Fact 7, f can be approximated from $O(\|F\|_4^4/\gamma\varepsilon^2)$ or $O(\|F\|_{\infty}^4 \log(1/\gamma)/\varepsilon^2)$
 359 queries to F in linear time. Thus, it admits an efficient approximation oracle (i.e., $\|F\|_4^4$ is bounded).

360 [5] reduces the problem of minimizing a submodular f to minimizing its Lovász extension
 361 $\mathbb{E}_{\lambda \sim [0,1]}[f(\{i: x_i \geq \lambda\})]$, which is a convex function over $[0, 1]^n$. Applying the ellipsoid method,
 362 they obtain an efficient algorithm for minimizing f up to arbitrary accuracy. Their best algorithm has
 363 complexity that is inverse logarithmic in the accuracy parameter but assumes the availability of an
 364 exact oracle. An approximate oracle for f can still emulate a *weak* separation oracle for the level sets
 365 of the Lovász extension, implying the our main POAD_SM algorithmic consequence.